

✓ Part 1: Regression Task

```
# Import the libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, root_mean_squared_error
from tabulate import tabulate
import random
random.seed(42)
np.random.seed(42)
```

```
# Exploring the data
df = pd.read_csv('Stopping.csv')
print(df)
```

```
↩
  Speed  Distance
0      4         4
1      5         2
2      5         4
3      5         8
4      5         8
..     ...      ...
57     35       107
58     36        79
59     39       138
60     40       110
61     40       134

[62 rows x 2 columns]
```

```
# Checking the data type:
df.dtypes
```

```
↩
      0
Speed  int64
Distance  int64

dtype: object
```

```
# Checking if there are any null values in the dataset
df.isnull().any()
```

```
↩
      0
Speed  False
Distance  False

dtype: bool
```

Task 1.1

```
# Linear Regression
# Declaring the variables
# X = independant variable: Speed (needs to be a dataframe, c so use [[]])
# y = dependant variable: Distance
X = df[['Speed']]
y = df['Distance']
```

```
X.shape
print(type(X))
# X = X.reshape(-1,1)
```

```
↩ <class 'pandas.core.frame.DataFrame'>
```

```
# splitting the dataset for training and testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30, random_state=42)
```

```
# building the Linear Regression model
#  $y = mX + b$ 
lr = LinearRegression()
```

```
# fitting the data in the model
lr.fit(X_train, y_train)
```

```
↗ LinearRegression ⓘ ?
LinearRegression()
```

```
# coefficient or slope of the line
# the lr.coef_ will tell us that an increase in speed will lead to an increase in distance by output amount
print(lr.coef_)
```

```
↗ [2.9667382]
```

```
# when speed = 0, distance = intercept
print(lr.intercept_)
```

```
↗ -18.879653658049712
```

```
# predicting distance based on speed
y_pred = lr.predict(X_test)
```

```
# printing the predicted values
print(y_pred)
print(y_test.values)
```

```
↗ [67.15575407 84.95618325 -7.01270087 84.95618325 1.88751372 64.18901587
 10.78772832 7.82099012 25.6214193 96.82313604 31.5548957 4.85425192
 61.22227767 7.82099012 34.5216339 37.48837209 -4.04596267 16.72120471
 70.12249226]
 [ 54  85   4 107   7  64  17   5  16 138  29  11  78  13  47  30   8  11
 101]
```

```
# building the Polynomial Regression model with degree = 2
#  $y = \beta + \beta_1X + \beta_2X^2$ 
pr = PolynomialFeatures(degree = 2)
X_train_p = pr.fit_transform(X_train)
X_test_p = pr.transform(X_test)
```

```
# training the model
poly_lr = LinearRegression()
poly_lr.fit(X_train_p, y_train)
```

```
↗ LinearRegression ⓘ ?
LinearRegression()
```

```
# predicting
y_pred_p = poly_lr.predict(X_test_p)
```

```
# Coefficient for polynomial regression with degree 2
print(poly_lr.coef_)
```

```
↗ [0.          0.63191779 0.05603184]
```

```
# intercept for polynomial regression with degree 2
print(poly_lr.intercept_)
```

```
↗ 0.3966086819856045
```

```
# Polynomial Regression model with degree = 3
#  $y = \beta + \beta_1X + \beta_2X^2 + \beta_3X^3$ 
pr3 = PolynomialFeatures(degree = 3)
X_train_p3 = pr3.fit_transform(X_train)
X_test_p3 = pr3.transform(X_test)
```

```
# training the model
poly_lr3 = LinearRegression()
poly_lr3.fit(X_train_p3, y_train)
```

LinearRegression

LinearRegression()

```
# predicting
y_pred_p3 = poly_lr3.predict(X_test_p3)

# Coefficient for polynomial regression with degree 3
print(poly_lr3.coef_)

[ 0.00000000e+00  1.98562408e+00 -1.45170070e-02  1.05882480e-03]

# intercept for polynomial regression with degree 3
print(poly_lr3.intercept_)

-6.571054534272392
```

Task 1.2

```
# to test on linear regression test dataset
# using tabulate for better readability
table = list(zip(X_test['Speed'], y_test, y_pred))
print(tabulate(table, headers = ["Speed", "Test distance", "Predicted distance"]))
```

Speed	Test distance	Predicted distance
29	54	67.1558
35	85	84.9562
4	4	-7.0127
35	107	84.9562
7	7	1.88751
28	64	64.189
10	17	10.7877
9	5	7.82099
15	16	25.6214
39	138	96.8231
17	29	31.5549
8	11	4.85425
27	78	61.2223
9	13	7.82099
18	47	34.5216
19	30	37.4884
5	8	-4.04596
12	11	16.7212
30	101	70.1225

```
# For Linear Regression
lr_MAE = mean_absolute_error(y_test, y_pred)
lr_MSE = mean_squared_error(y_test, y_pred)
lr_RMSE = root_mean_squared_error(y_test, y_pred)
lr_R2 = r2_score(y_test, y_pred)
print(f" Linear Regression MAE = {lr_MAE}")
print(f" Linear Regression MSE = {lr_MSE}")
print(f" Linear Regression RMSE = {lr_RMSE}")
print(f" Linear Regression R2 = {lr_R2}")
```

```
Linear Regression MAE = 11.087259207505738
Linear Regression MSE = 228.25103406651817
Linear Regression RMSE = 15.10797915230618
Linear Regression R2 = 0.8566481849373467
```

```
# For Polynomial Regression of degree 2
pr_MAE = mean_absolute_error(y_test, y_pred_p)
pr_MSE = mean_squared_error(y_test, y_pred_p)
pr_RMSE = root_mean_squared_error(y_test, y_pred_p)
pr_R2 = r2_score(y_test, y_pred_p)
print(f" Polynomial Regression (Degree = 2) MAE = {pr_MAE}")
print(f" Polynomial Regression (Degree = 2) MSE = {pr_MSE}")
print(f" Polynomial Regression (Degree = 2) RMSE = {pr_RMSE}")
print(f" Polynomial Regression (Degree = 2) R2 = {pr_R2}")
```

```
Polynomial Regression (Degree = 2) MAE = 8.726619840934255
Polynomial Regression (Degree = 2) MSE = 158.21962845767766
Polynomial Regression (Degree = 2) RMSE = 12.578538407051816
Polynomial Regression (Degree = 2) R2 = 0.9006310266645413
```

```
# For Polynomial Regression of degree 3
pr3_MAE = mean_absolute_error(y_test, y_pred_p3)
pr3_MSE = mean_squared_error(y_test, y_pred_p3)
```

```
pr3_RMSE = root_mean_squared_error(y_test, y_pred_p3)
pr3_R2 = r2_score(y_test, y_pred_p3)
print(f" Polynomial Regression (Degree = 3) MAE = {pr3_MAE}")
print(f" Polynomial Regression (Degree = 3) MSE = {pr3_MSE}")
print(f" Polynomial Regression (Degree = 3) RMSE = {pr3_RMSE}")
print(f" Polynomial Regression (Degree = 3) R2 = {pr3_R2}")
```

```
Polynomial Regression (Degree = 3) MAE = 9.002675448876635
Polynomial Regression (Degree = 3) MSE = 161.4531663131979
Polynomial Regression (Degree = 3) RMSE = 12.706422246769462
Polynomial Regression (Degree = 3) R2 = 0.8986002208784544
```

The above methods we have used to evaluate the performance of regression models includes us calculating various metrics upon 2 values -

1. Authentic values (y_{test})
2. Anticipated values ($y_{\text{pred}}/y_{\text{pred_p}}/y_{\text{pred_p3}}$)

The difference or similarity between these 2 sets tells us all about the performance. We have made use of scikit-learn which is a highly regarded and widely used machine learning library. The functions we will be using are as follows:

1. Mean absolute error (MAE) - This will tell us about the average difference between authentic and anticipated values. For example, if the MAE is 6.2 then the error between expected speed and actual speed on an average is 6.2 km
2. Mean squared error (MSE) - Squared difference between authentic and anticipated values. The error value is squared before it is averaged. This helps in understanding the error occurring on larger values.
3. Root mean squared error (RMSE) - Average difference between authentic and anticipated values after being squared. Basically the square root of MSE.
4. R^2 score (R^2) - Tells us how well the model fits the data (variance). This value lies between 0 and 1. If the model has a R^2 score of 96%, it means that the model can explain 96% of variation in distance according to speed.

Task 1.3

Underfitting is when a model doesn't capture all the pattern in the data or is too simple. Underfit models have a low variance and high error.

Overfit models are the opposite of underfit models i.e. they capture everything in the data along with additional unnecessary data. Overfit models give a perfect score during training but work poorly on testing/unseen data.

According to our findings -

For linear regression:

MAE is the highest when compared to other models, which tells us that the overall prediction of distance is off by ~81 feet.

MSE and RMSE of linear regression model are almost double when compared to other models, so the model has made significant errors while predicting large stopping distances.

R^2 test tells us that the model is able to explain 86% of the variance in stopping distance based on speed.

This model exhibits high errors as seen after calculating MAE, MSE, and RMSE which can be a case of underfitting. Since it's a linear regression model, it assumes the relationship to be a straight line and possibly won't give us accurate results by considering the complexity of the data in real life.

For polynomial regression (degree = 2):

MAE tells us that the overall prediction is off only by ~8.7 feet which is a large improvement when compared to our linear regression model. The lower the MAE, the better the model.

MSE and RMSE are much lower than that of the linear regression model. This means this model is somewhat reliable at predicting the distance at large stopping distances.

R^2 test explains 95% of variance in stopping distance based on speed.

Our observations have improved immensely as observed by comparing the MAE, MSE, and RMSE with the linear regression model. This model introduces curvature which captures the actual relationship between speed and distance. Since the errors observed are smaller and the model is able to explain a high amount of variance, we can say that this model looks like a perfect fit for our data up until now.

For polynomial regression (degree = 3):

MAE for this model has exhibited a further reduction in error prediction compared to the degree 2 polynomial regression model, with this model's MAE being ~6.2.

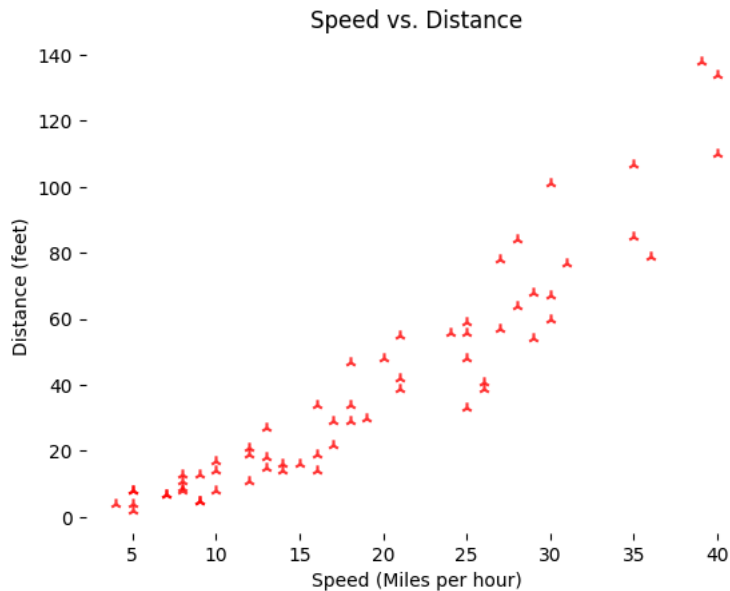
MSE and RMSE were also slightly lower than the previous model, making polynomial regression with degree 3 the best fit model up till now.

A slight improvement in R^2 test to 96%, indicating an even better explanation of variance compared to polynomial regression (degree = 2).


The R^2 score has increased slightly and the model's calculated error metrics are better than both the linear regression and polynomial regression (degree 2) models. This means that while the model has increased in complexity, it translates to better performance and not overfitting in this case. Final verdict is that polynomial regression degree 3 is the best performing model overall.

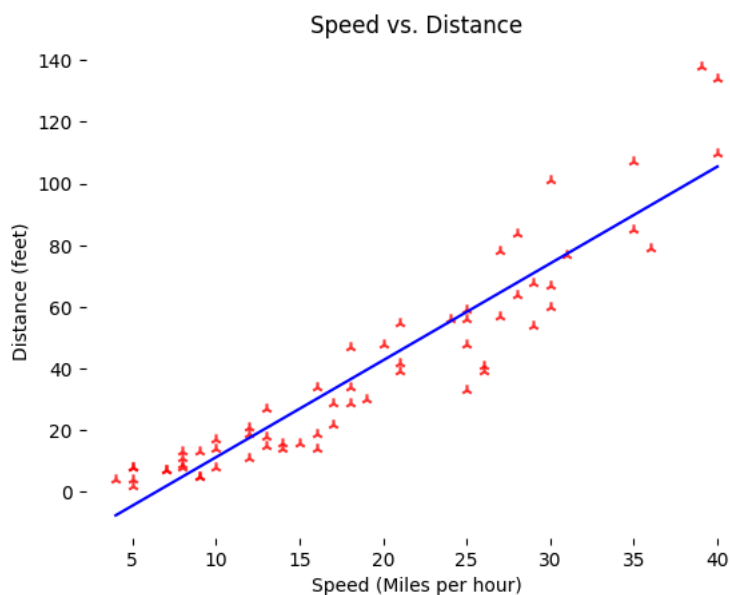
Task 1.4

```
# Visualising the dataset with the help of scatterplot
X = df.Speed
y = df.Distance
plt.scatter(X,y, color = 'red', marker = "2", alpha = 0.8)
plt.title("Speed vs. Distance")
plt.xlabel('Speed (Miles per hour)')
plt.ylabel('Distance (feet)')
plt.box(False)
plt.show()
```



```
# Visualising linear regression plot
lr = LinearRegression()
lr.fit(pd.DataFrame(X), y)
pred = lr.predict(X.values.reshape(-1,1))
plt.scatter(X, y, color = 'red', marker = "2", alpha = 0.8)
plt.plot(X, pred, color = 'blue')
plt.title("Speed vs. Distance")
plt.xlabel('Speed (Miles per hour)')
plt.ylabel('Distance (feet)')
plt.box(False)
plt.show()
```

 /usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning: X does not have valid feature names, but LinearRegression was fitted to X
warnings.warn()




```
# Visualising polynomial regression of degree 2 plot
pr = PolynomialFeatures(degree = 2)
pX = pr.fit_transform(X.values.reshape(-1,1))
lr = LinearRegression()
```

```

lr.fit(pX, y)
pred = lr.predict(pX)
plt.scatter(X, y, color = 'purple', marker = "2", alpha = 0.8)
plt.plot(X, pred, color = 'orange')
plt.title("Speed vs. Distance")
plt.xlabel('Speed (Miles per hour)')
plt.ylabel('Distance (feet)')
plt.box(False)
plt.show()


```



```

# Visualising polynomial regression of degree 3 plot
pr = PolynomialFeatures(degree = 3)
pX = pr.fit_transform(X.values.reshape(-1,1))
lr = LinearRegression()
lr.fit(pX, y)
pred = lr.predict(pX)
plt.scatter(X, y, color = 'green', marker = "2", alpha = 0.8)
plt.plot(X, pred, color = 'maroon')
plt.title("Speed vs. Distance")
plt.xlabel('Speed (Miles per hour)')
plt.ylabel('Distance (feet)')
plt.box(False)
plt.show()

```



```

# Visualising all the plots together
X_resaped = X.values.reshape(-1, 1)

```

```

# Linear Regression
lr = LinearRegression()
lr.fit(pd.DataFrame(X), y)
pred_lr = lr.predict(X_resaped)

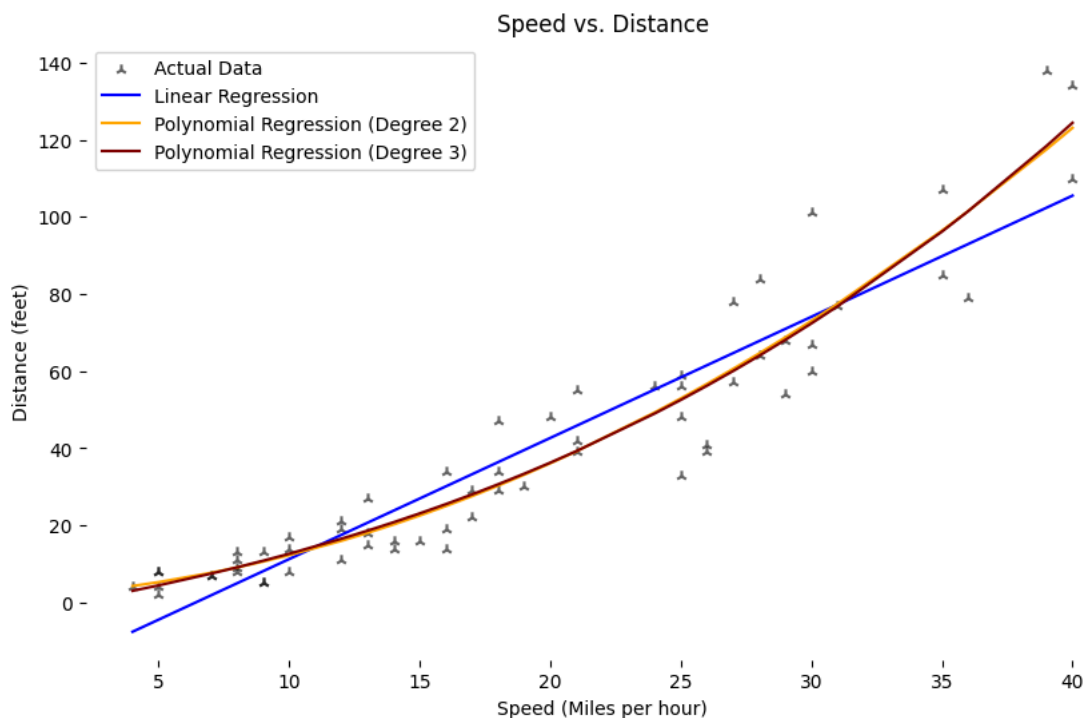
```

```
# Polynomial Regression (degree = 2)
pr2 = PolynomialFeatures(degree=2)
pX2 = pr2.fit_transform(X_resaped)
lr2 = LinearRegression()
lr2.fit(pX2, y)
pred_p2 = lr2.predict(pX2)

# Polynomial Regression (degree = 3)
pr3 = PolynomialFeatures(degree=3)
pX3 = pr3.fit_transform(X_resaped)
lr3 = LinearRegression()
lr3.fit(pX3, y)
pred_p3 = lr3.predict(pX3)

plt.figure(figsize=(10,6))
plt.scatter(X, y, color='black', marker="2", alpha=0.6, label='Actual Data')
plt.plot(X, pred_lr, color='blue', label='Linear Regression')
plt.plot(X, pred_p2, color='orange', label='Polynomial Regression (Degree 2)')
plt.plot(X, pred_p3, color='maroon', label='Polynomial Regression (Degree 3)')
plt.title("Speed vs. Distance")
plt.xlabel('Speed (Miles per hour)')
plt.ylabel('Distance (feet)')
plt.legend()
plt.box(False)
plt.show()
```

🔗 /usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning: X does not have valid feature names, but LinearRegression was fitted to X
warnings.warn()



✓ Part 2: Classification Task

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, roc_curve, auc
import random
random.seed(42)
np.random.seed(42)
```

```
# Loading and displaying the dataset
df = pd.read_csv('CancerClassification.csv')
print(df)
```

🔗

	Cl.thickness	Cell.size	Cell.shape	Marg.adhesion	Epith.c.size	\
0	45.112062	51.798322	47.877500	53.385136	58.408301	
1	70.864529	57.307064	63.333871	62.865781	55.409522	

```

2      49.307810  71.424268  42.062140    40.255089   35.533395
3      66.775425  60.810981  64.712988    79.132088   67.626699
4      52.012773  62.644698  67.966587    43.309239   51.623986
..      ...      ...      ...      ...      ...
678    45.895448  58.859543  63.878934    59.358741   50.247380
679    68.697373  37.409173  49.989872    49.115574   55.864530
680    51.550280  69.778100  54.258628    55.536820   57.292757
681    63.460881  49.141312  38.042230    61.115623   33.504183
682    51.766514  41.607881  64.059539    62.313476   44.386715

```

```

      Bare.nuclei  Bl.cromatin  Normal.nucleoli  Mitoses  Class
0      39.929594   46.499295    50.669261   53.078398     0
1      41.867581   55.753702    52.593935   54.329550     0
2      48.339003   65.823549    46.045521   37.110279     0
3      54.239235   63.276384    42.660030   73.929424     0
4      43.002597   70.405915    43.635677   55.973108     0
..      ...      ...      ...      ...      ...
678    67.011975   56.875841    64.163137   53.106369     0
679    57.494473   64.858096    40.635358   56.281143     0
680    54.082691   55.668014    58.590465   76.758074     1
681    71.588486   68.387738    65.392520   38.025973     1
682    58.153828   69.110012    34.321386   58.935607     1

```

[683 rows x 10 columns]

Exploring the data

df.columns

```

Index(['Cl.thickness', 'Cell.size', 'Cell.shape', 'Marg.adhesion',
      'Epith.c.size', 'Bare.nuclei', 'Bl.cromatin', 'Normal.nucleoli',
      'Mitoses', 'Class'],
      dtype='object')

```

df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 683 entries, 0 to 682
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Cl.thickness         683 non-null   float64
1   Cell.size            683 non-null   float64
2   Cell.shape          683 non-null   float64
3   Marg.adhesion        683 non-null   float64
4   Epith.c.size         683 non-null   float64
5   Bare.nuclei          683 non-null   float64
6   Bl.cromatin          683 non-null   float64
7   Normal.nucleoli      683 non-null   float64
8   Mitoses              683 non-null   float64
9   Class                683 non-null   int64
dtypes: float64(9), int64(1)
memory usage: 53.5 KB

```

df.describe()

```

      Cl.thickness  Cell.size  Cell.shape  Marg.adhesion  Epith.c.size  Bare.nuclei  Bl.cromatin  Normal.nucleoli  Mitoses
count      683.000000    683.000000    683.000000    683.000000    683.000000    683.000000    683.000000    683.000000    683.000000
mean        53.794592    54.493624    54.902298     53.704678     53.722228     54.141138     53.836756     54.427693     53.636756
std         10.332521    10.368235     9.999999     10.669766     10.474714     10.146594     10.756423     10.102067     10.927999
min         11.641612    23.947718    24.821445     20.475761     23.598114     19.397066     21.727114     23.159512     24.521250
25%         46.575414    47.319452    47.773908     45.897371     46.664622     47.945944     47.385312     47.882020     46.578250
50%         54.239844    54.508344    55.160264     53.432298     53.815078     53.823819     54.120836     54.539211     54.060250
75%         60.300560    61.436500    62.065397     60.865782     60.741584     61.112674     60.397951     60.919610     60.357350
max         89.699515    93.321471    87.376942     85.833716     89.550155     85.626194     79.255614     85.948280     86.227400

```

df.isnull().sum()

	0
Cl.thickness	0
Cell.size	0
Cell.shape	0
Marg.adhesion	0
Epith.c.size	0
Bare.nuclei	0
Bl.cromatin	0
Normal.nucleoli	0
Mitoses	0
Class	0

dtype: int64


```
# Exploring our target
df['Class'].value_counts()
```

	count
Class	
0	444
1	239

dtype: int64

```
# Visualizing our target using sns plot
sns.countplot(x='Class', data=df)
```

<Axes: xlabel='Class', ylabel='count'>



Task 2.1

```
# Separating the data into features and target (X and y)
X = df.drop('Class', axis = 1)
y = df['Class']
```

```
# Train test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Selecting and fitting the model
model = LogisticRegression()
model.fit(X_train, y_train)
```

LogisticRegression

```
y_pred = model.predict(X_test)
```

```
# First 10 predictions
print(y_pred[:10])
```

```
# Actual values
print(y_test.values[:10])
```

```
↩ [0 0 1 0 0 0 0 0 0 0]
   [1 1 0 0 0 0 0 0 1 0]
```

Task 2.2

```
print(confusion_matrix(y_test, y_pred))
```

```
↩ [[72  7]
   [36 22]]
```

```
print(classification_report(y_test, y_pred))
```

```
↩
```

	precision	recall	f1-score	support
0	0.67	0.91	0.77	79
1	0.76	0.38	0.51	58
accuracy			0.69	137
macro avg	0.71	0.65	0.64	137
weighted avg	0.71	0.69	0.66	137

```
score = accuracy_score(y_test, y_pred)
print(score)
```

```
↩ 0.6861313868613139
```

Confusion Matrix is a method which helps us in summarising the performance metrics of our model. We understand the specificity, sensitivity and accuracy of our model with the help of the confusion matrix. We have a 2x2 matrix with each cell representing the comparison between the prediction and actual values.

In cell 1x1, which is True Negative (TN) – the model predicts 0 while the actual value is 0. The model was successful 72 times.

In cell 1x2: False Positive (FP) – the model predicts 0 as 1; 7 times.

In cell 2x1: False Negative (FN) – the model predicts 1 as 0; 36 times.

In cell 2x2: True Positive (TP) – the model predicts 1 as 1; 22 times.

The classification report gives us precision, recall and F1-score of our model in a tabular format.

Precision gives us the amount of correct positive predictions and the formula is: $\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$

Recall gives us the proportion of actual positives our model was able to identify and it is given by: $\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$

F1-score gives us a balance between Precision and Recall, and its formula is: $\text{F1-score} = 2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$

Accuracy Score gives us a percentage of the number of accurate predictions made by our model, divided by the total number of predictions:

$\text{Accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN})$


Task 2.3

The overall accuracy of our model is 68.6%, a decent score but considering the context of our project this cannot be the only defining metric. Addressing the output of our classification report reveals several insights regarding the performance of our model. The model's precision for malignant (Class 1) is 0.76, meaning when it predicts a case as malignant, it is correct 76% of the time. Recall for our model in Class 0 is 0.91 and Class 1 is 0.38; this is concerning since our model is far better at identifying cases where cancer cells are benign than cases where cancer cells are malignant. Ideally, recall for malignant (Class 1) cases should be as high as possible, since it is very important to not miss an actual cancer diagnosis, but our model is only able to correctly identify 38% of actual malignant cases. F1-score for Class 0 of this model is 0.77, which is good, but the model should perform better at predicting Class 1 which has an F1-score of only 0.51. Support shows our test data has 79 cases of Class 0 which is slightly higher than the 58 cases of Class 1, indicating a mild class imbalance.


The most important metric in our context of breast cancer classification is Recall for Class 1, which refers to malignant cases, since it is essential to identify patients who truly have cancer. Here, false negatives can lead to fatal consequences. Since we have a greater number of Class 0, accuracy will be ineffective as the model can just predict "Benign" most of the time and have a similar performance. This is called class imbalance and it affects the performance of our model. Precision for Class 1 is also good, indicating that when the model identifies a case as malignant, it is correct most of the time. However, it is less critical than recall in our context.

Task 2.4

```
# confusion matrix heatmap
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt="d")
```




```
# classification report barplot
cr = classification_report(y_test, y_pred, output_dict=True)
df_cr = pd.DataFrame(cr).transpose()[['precision', 'recall', 'f1-score']]
df_cr.plot(kind='bar', color = ['slateblue', 'mediumseagreen', 'tomato'])
plt.title('Classification Report (Precision, Recall, F1-score)')
plt.ylabel('Score')
plt.ylim(0, 1)
plt.xticks(rotation = 0)
plt.tight_layout()
plt.grid(axis = 'y')
plt.show()
```




```
# Class imbalance
df['Class'].value_counts().plot(kind='pie', autopct='%1.0f%%', labels = ['Class 0', 'Class 1'], colors = ['slateblue', 'tomato'])
plt.title('Class Distribution (Class 0 vs Class 1)')
plt.show()
```



Class Distribution (Class 0 vs Class 1)



```
# ROC and AUC
y_pred_prob = model.predict_proba(X_test)[: , 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
roc_auc= auc(fpr, tpr)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



Start coding or [generate](#) with AI.